

Stat 8120 Final Project

Joe Paul*

Introduction

In this project I attempted to generate a model for TNC fare from the publicly available Chicago TNC dataset for 2023+.¹ In my research, I work to improve the usage of pooled rideshare. Pooled rideshare is a service in which you share part of a rideshare trip with a stranger. Uber and Lyft offer such services, namely Uber Pool and Lyft Shared. In this project, led by Dr. Yunyi Jia and Dr. Johnelle Brooks, approached the topic of pooled rideshare from several different angles. Dr. Brooks lab focuses on the human factors of ridesharing and work with JD Power to construct national surveys designed to further the understanding of factors that impact people use of such services. Through two national surveys, we have learned that safety is the most important factor that prevents usage of pooled rideshare. More results from the analysis of the survey data can be viewed in² and.³ This data was also utilized to create a choice model which helps predict whether a given rider will choose a pooled ride when presented with a choice between a pooled and solo ride.⁴

In addition to modeling and human factors work, we also tackle the ridership problem from the angle of transportation network company (TNC) operators, like Uber and Lyft. This section is my responsibility, and I attempt to address these issues through designing new fleet control algorithms that improve the service experienced by the riders. Given that TNC fleets are typically composed of contracted drivers utilizing their own vehicles, the best way to address level of service (LOS) issues like long waits or delays associated with pooled rides, is to control the fleets operations through fleet control algorithms. Fleet control can be divided into three primary categories as illustrated below in 1.

My work thus far has been to design an assignment algorithm and to implement and test it in a traffic simulation application called POLARIS.⁵ The assignment algorithm, detailed in my paper,⁶ utilizes a willingness to pay formulation of the pooling choice model detailed in⁴ to estimate the discount required to incentivize a rider to pool. The primary incentive for riders in pooled rideshare to pool is the cost savings over a solo ride. One important aspect to this study is to understand the base rate that riders should be charged. This is important because it helps us understand whether or not a system is likely to be profitable and provides a basis from which we can discount a trip. To estimate fares, a model is required.

I have previously generated a simple LM from,¹ but wanted to take this opportunity and the new skills learned in this class to rehash the model. Additionally, my original model was calibrated from the TNC data ending in Fall of 2022, and as inflation has heavily affected the US in the last

Algorithm	Explanation
Assignment	Control of which vehicles are matched to which request, including matching multiple requests to single vehicles
Repositioning	Moving vehicles that are not currently servicing requests to more optimal locations for servicing expected future demand
Routing	Controlling the turn by turn directions that vehicles use to navigate the traffic network

Table 1: Fleet control algorithms

several years, I felt it was necessary to recalibrate using the newest available data. Additionally, as several other factors are available in this dataset, and my original model only utilized the time and distance for its estimates, I wanted to use the feature selection techniques learned to determine whether my uninformed selection was satisfactory. To that end, the goal of this project is to calibrate a new fare model from the newest available TNC data and to utilize feature selection and advanced model generation techniques to find the best model for fare.

Outline of Analysis

To address this task, I will first draw a sample from the Chicago TNC dataset. I will then utilize some basic exploratory analysis including GG plots and summary statistics to understand which variables contribute to fare, and to identify any problem variable that should not be utilized in the model. I will start with a simple LM utilizing the significant variables. Since this dataset contains relative few predictors, especially relative to the number of samples, no dimension reduction should be necessary. However, I will utilize subset regression, stepwise selection with cross validation, and principal component analysis to help narrow in on the best model structure and most important variables. Finally, I will fit the final model and interpret its results.

Methodology

To calibrate a model, I first obtained a copy of the most recently available data through the City of Chicago's web data portal. I filtered the data for June only (the latest available date) to reduce the output size. Still, the dataset generated an over 2GB CSV containing 750,000 individual

*jpaul4@clemson.edu

trips. The columns are illustrated in 2.

Next I utilized sqlite to convert the CSV into a sqlite database to pull a sample of rows small enough to work with. I used the "random" function in sqite to capture a random sample of trips accross the month. I pulled 2000 trips in total to use and split it into a test and train dataset with a 90/10 split. Prior to doing this, I used the sqldf package to query the 10 first rows to look at the data.

Next, I was curious whether the day of week had any affect on the fares, so I generated a day-of-week (DOW) column from the trip start timestamp column. Next, I removed any columns that I knew were not of interest, like the geographic columns and ID columns. Given that the total cost column was supposedly a sum of the three cost related variable columns, I decided to check if it matched the sum. In some cases it did not, and given its dependency on these columns anyways, I discarded it to remove colinearity. I then checked the rows within the sample for NA's in any of the remaining columns. Two had NAs so I removed these two rows.

Next, I began the work for model fitting by separating the dataframe into a train and test dataset in a 90/10 mix. I used the GGpairs function to inspect the relationship between the variables remaining in the training set. The GG pairs result is illustrated in 1.

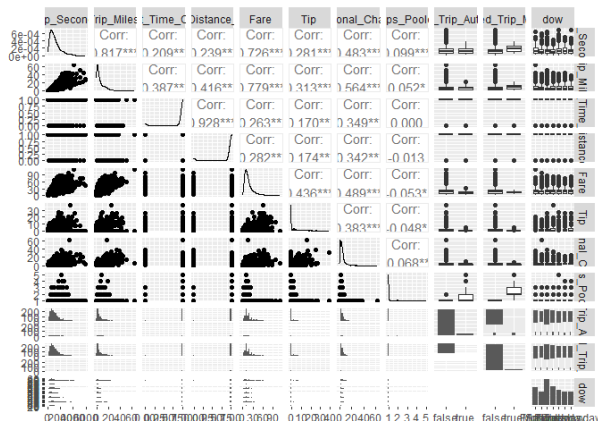


Figure 1: GGpairs results for training dataset.

While inspecting the GG pairs columns, I noticed that the distance and time in chicago columns were mostly 0, so I removed these from the dataframe also. They also would inhibit the models scalability to other regions. I removed them from the dataframe and reformed by test and train sets on the same test indices to ease the process. First, I used the LM function to inspect the significance of the remaining variables to predicting fare. I obtained the following output as illustrated in 3. The LM function automatically converted categorical variables to binary columns.

Analyzing the results, I found that trip-seconds, trip-miles, tip, and shared-trip-authorized were significant at a 0 level , and that dow sunday was significant at a 0.1 level. Since intuitively, tip was likely very colinear with fare, I formed a correlation matrix. Trip minutes and trip miles were highly

Variable	Description
PK-UID	Primary Index
Trip-ID	Coded trip identifier
Trip-Start-Timestamp	Time/Date pickup rounded to nearest 15 mins
Trip-End-Timestamp	time/Date dropoff rounded to nearest 15 mins
Trip-Seconds	trip time in seconds
Trip-Miles	trip distance in miles
Percent-Time-Chicago	percent time in Chicago city
Percent-Distance-Chicago	percent distance in Chicago
Pickup-Census-Tract	location reference to census data
Dropoff-Census-Tract	location reference to census data
Pickup-Community-Area	location reference to area name (encoded)
Dropoff-Community-Area	location reference to area name (encoded)
Fare	fare in dollars rounded to the nearest \$2.50
Tip	Tip paid to driver rounded to the nearest \$1
Additional-Charges	taxes, fees, other miscellaneous charges
Trip-Total	sum of tip, fare, and additional charges including rounding
Shared-Trip-Authorized	boolean for if a trip was accepted to be pooled regardless of whether a match was found
Shared-Trip-Match	boolean if match was found
Trips-Pooled	if a match was found, how many (including the base trip)
Pickup-Centroid-Latitude	geographic location of pickup tract centroid
Pickup-Centroid-Longitude	geographic location of pickup tract centroid
Pickup-Centroid-Location	geographic location of pickup tract centroid (stored in geo data form)
Dropoff-Centroid-Latitude	geographic location of pickup tract centroid
Dropoff-Centroid-Longitude	geographic location of pickup tract centroid
Dropoff-Centroid-Location	geographic location of pickup tract centroid (stored in geo data form)

Table 2: Columns in TNC Trip Dataset

Coefficient	Estimate	StdError	tvalue	Pr(> t)
(Intercept)	6.894924	1.621669	4.252	2.23e-05 ***
Trip Seconds	0.004439	0.000379	11.714	<2e-16 ***
Trip Miles	0.930867	0.045233	20.579	<2e-16 ***
Tip	0.798935	0.059870	13.344	<2e-16 ***
Additional Charges	-0.030049	0.053778	-0.559	0.5764
Trips Pooled	-0.792071	1.531302	-0.517	0.6050
Shared Trip Authorized true	-4.890068	1.067234	-4.582	4.92e-06 ***
Shared Trip Match true	-2.971141	2.640581	-1.125	0.2607
dow Monday	-0.085948	0.729787	-0.118	0.9063
dow Saturday	0.458279	0.623896	0.735	0.4627
dow Sunday	1.223020	0.671381	1.822	0.0687 .
dow Thursday	-0.705050	0.703864	-1.002	0.3166
dow Tuesday	-0.613870	0.728873	-0.842	0.3998
dow Wednesday	-0.520356	0.725991	-0.717	0.4736

Table 3: Coefficients and significances.

correlated, and tip and additional charges were moderately correlated. I decided to remove tip knowing that it was likely a function of fare, and not useful predictor since it would not be known at the start of a trip anyways. I did the same with additional charges, match found, and trips pooled as they are also not likely to be known prior to a trip commencing. I reran the LM to check if any significances had been affected, and now dow Sunday was significant at the 0.05 level. The resulting LM, for minutes, miles, dow, and share trip authorized had an r^2 value of 0.6512 with an MSE when used to predict on the test dataset of 52.22.

Next, I decided to continue my exploratory analysis by using subset regression. Of the 8 remaining variables, the subset of three variables illustrated the biggest transition across the evaluation metrics RSS, adjusted r^2 , Cp and BIC. The charts for these at each subset are illustrated in 2.

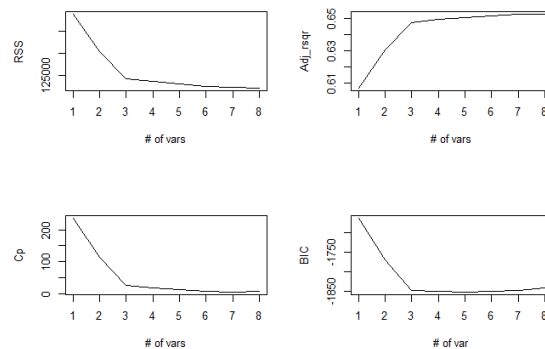


Figure 2: RSS, adjusted r^2 , Cp and BIC for subset regression.

Although some of the higher order models showed improvement, it was apparent that the 3 variable model was the best fit due to its simplicity and transition point in the evaluative metrics. The 3 variable model utilized the trip seconds, trip miles, and shared trip authorized variables. Intuitively, this is what I was expecting, but, it was helpful to validate that DOW did not contribute significantly.

Please note, that set.seed does not appear to work properly in r markdown documents, and that these values may change in the final output r markdown document due to this fact. These results were written on a complete set of outputs, but if the workbook is reprocessed, they may change as the training sample will change.

Next, I utilized cross validation stepwise selection. The mean cv errors in this case yielded the lowest error in the 7 variable model, but, the 3 variable model look to be the best due to the spike in cv error at the 4 variable model. The mean cv error plot is illustrated in 3.

The model generated from the stepwise selection process yielded an MSE of 51.644 with an r^2 of 0.63. This is slightly worse than the r^2 from the LM but the MSE is slightly better.




Figure 3: Mean CV Errors from Stepwise Selection




Figure 4: PLS Principal Components MSEP




Figure 5: PLS automatic feature selection RMSEP.

Finally, principle component analysis was utilized. The model with 11 principal components had the best MSEP, but the 2 principal component model illustrated the largest drop, and therefore is the best selection. The PLS principal component MSEP errors are illustrated in 4 and the automatic feature selection function, "selectNcomp" was used to generate an additional chart with RMSEP as show in 5. The model generated from the PLS model yielded and MSE of 52.56, slightly higher than the LM.

Conclusions

Given the simplicity, interpretability, and predictive accuracy of the LM, I selected it as the final model. I selected the shared trip authorized, trip miles, and trip seconds variables as they were the most clearly significant. I fit a final model with just these three variables, and obtained the following coefficients as illustrated in 4.

Variable	Value
Intercept	\$6.60
Shared Trip Authorized (true = 1)	\$-7.80
Trip Miles	\$1.02 per mile
Trip Seconds	\$0.004737 (\$0.27 per minute)
MSE	51.30
r^2	0.648

Table 4: Final Fare Model

Removing the DOW improved the predictive accuracy slightly, making this the best performing model tested. The values obtained for each variable are reasonable and interpretable. To utilize the model for prediction, eqn. 1 illustrates the complete formulation.

$$f = 6.6 + -7.8 * s_a + 1.02 * t_m + 0.27 * t_d \quad (1)$$

s_a is 1 if pooling is authorized, and zero otherwise, t_m is the trip distance in miles and t_d is the trip length in minutes.

through this study, a variety of exploratory and explanatory methods were utilized in pursuit of constructing an accurate TNC fare model for use in my ridesharing research. Although the initial LM was nearly the best model, the other methods utilized in this study helped illustrate and confirm that the simple LP was sufficient to capture the best model.

Appendix

Please see R Markdown below.

Setup

```
library(data.table)
library(sqldf)

## Loading required package: gsubfn
## Loading required package: proto
## Loading required package: RSQLite

library(ggplot2)
library(GGally)

## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2

library(leaps)
library(MASS)
library(glmnet)

## Loading required package: Matrix
## Loaded glmnet 4.1-8

library(pls)

##
## Attaching package: 'pls'

## The following object is masked from 'package:stats':
##
##   loadings

library(boot)
library(gridExtra)
library(splines)
library(caret)

## Loading required package: lattice

##
## Attaching package: 'lattice'

## The following object is masked from 'package:boot':
##
##   melanoma

##
## Attaching package: 'caret'

## The following object is masked from 'package:pls':
##
##   R2

library(alr4)

## Loading required package: car
## Loading required package: carData

##
## Attaching package: 'car'
```

```
## The following object is masked from 'package:boot':
##
##   logit

## Loading required package: effects

## Use the command
##   lattice::trellis.par.set(effectsTheme())
##   to customize lattice options for effects plots.
## See ?effectsTheme for details.
```

Read Data

This data csv is 2GB, way too big to read in on its own, so lets filter it down a bit.

```
csv_pth = "C:\\Users\\jpaul4\\Downloads\\Transportation_Network_Providers_-_Trips__2023-__20240801.csv"

#Get coloumn names
top_10 = sqldf::read.csv.sql(csv_pth, sql= "select * from file limit 10")
```

Select a sample from this CSV to use as our training data

I selected a random set of 2000 rows using the SQL command "SELECT * FROM chic_trips ORDER BY RANDOM() LIMIT 2000;" for training and testing.

```
csv_sample = "C:\\Users\\jpaul4\\Box\\Summer 2024\\S2405\\Projects\\Final Project\\result_set.csv"
trips = read.csv(csv_sample)
head(trips)
```

```
##      PK_UID      Trip_ID  Trip_Start_Timestamp
## 1 1078724 a45e404c67637a35ba8c68eb02abd0ef57a27851 06/26/2024 10:15:00 PM
## 2 5948098 9bd17c2ea621b7c1bba3725d89c18b45ac574971 06/07/2024 03:15:00 PM
## 3 1058729 92cc46f8531e33eee0e5a21cda3cb34d7997e0c8 06/27/2024 12:45:00 AM
## 4 2375826 4c4659ad8d22f0f0177e79e2140023a9ffea9ea2 06/21/2024 08:30:00 PM
## 5 559953 aceb9a135afb3a822bce97dbf5f80b31fa951c3d 06/28/2024 10:30:00 PM
## 6 7199958 8fd92610f23c8513b2618355ea3dbaf1ac734e37 06/02/2024 05:30:00 AM
##      Trip_End_Timestamp Trip_Seconds Trip_Miles Percent_Time_Chicago
## 1 06/26/2024 10:30:00 PM      1124      4.9      1
## 2 06/07/2024 03:30:00 PM      896      3.0      1
## 3 06/27/2024 01:30:00 AM     1994     27.8      0
## 4 06/21/2024 09:15:00 PM     2963     26.9      1
## 5 06/28/2024 10:45:00 PM      816      2.4      1
## 6 06/02/2024 06:00:00 AM     2260     32.4      1
##      Percent_Distance_Chicago Pickup_Census_Tract Dropoff_Census_Tract
## 1      1      NA      NA
## 2      1      NA      NA
## 3      0     17031980000      NA
## 4      1      NA      NA
## 5      1     17031081403     17031842200
## 6      1     17031820901     17031980000
##      Pickup_Community_Area Dropoff_Community_Area Fare Tip Additional_Charges
## 1      28      6 32.5  0      1.23
## 2      28     24 17.5  0      5.68
## 3      76     NA 60.0 12     21.39
## 4      NA      6 22.5  0      6.95
## 5       8      8 10.0  0      1.23
## 6      NA     76 47.5 12     28.12
##      Trip_Total Shared_Trip_Authorized Shared_Trip_Match Trips_Pooled
## 1      33.73      false      false      1
## 2      23.18      false      false      1
## 3      93.39      false      false      1
## 4      29.45      false      false      1
```

```

## 5      11.23                false                false                1
## 6      87.62                false                false                1
## Pickup_Centroid_Latitude Pickup_Centroid_Longitude
## 1              41.87400                -87.66352
## 2              41.87400                -87.66352
## 3              41.97907                -87.90304
## 4              NA                    NA
## 5              41.89092                -87.61887
## 6              NA                    NA
## Pickup_Centroid_Location Dropoff_Centroid_Latitude
## 1 POINT (-87.6635175498 41.874005383)                41.94423
## 2 POINT (-87.6635175498 41.874005383)                41.90121
## 3 POINT (-87.9030396611 41.9790708201)                NA
## 4              41.94423
## 5 POINT (-87.6188683546 41.8909220259)                41.90494
## 6              41.97907
## Dropoff_Centroid_Longitude Dropoff_Centroid_Location
## 1              -87.65600 POINT (-87.6559981815 41.9442266014)
## 2              -87.67636 POINT (-87.6763559892 41.9012069941)
## 3              NA
## 4              -87.65600 POINT (-87.6559981815 41.9442266014)
## 5              -87.64991 POINT (-87.6499072264 41.9049353016)
## 6              -87.90304 POINT (-87.9030396611 41.9790708201)

```

```
summary(trips)
```

```

## PK_UID Trip_ID Trip_Start_Timestamp Trip_End_Timestamp
## Min. : 2180 Length:2000 Length:2000 Length:2000
## 1st Qu.:1857128 Class :character Class :character Class :character
## Median :3728684 Mode :character Mode :character Mode :character
## Mean :3749083
## 3rd Qu.:5638863
## Max. :7588816
##
## Trip_Seconds Trip_Miles Percent_Time_Chicago
## Min. : 55.0 Min. : 0.200 Min. :0.0000
## 1st Qu.: 615.8 1st Qu.: 2.175 1st Qu.:1.0000
## Median : 991.5 Median : 4.400 Median :1.0000
## Mean :1218.3 Mean : 7.336 Mean :0.9325
## 3rd Qu.:1583.0 3rd Qu.:10.000 3rd Qu.:1.0000
## Max. :5937.0 Max. :63.900 Max. :1.0000
## NA's :1
## Percent_Distance_Chicago Pickup_Census_Tract Dropoff_Census_Tract
## Min. :0.0000 Min. :1.703e+10 Min. :1.703e+10
## 1st Qu.:1.0000 1st Qu.:1.703e+10 1st Qu.:1.703e+10
## Median :1.0000 Median :1.703e+10 Median :1.703e+10
## Mean :0.9275 Mean :1.703e+10 Mean :1.703e+10
## 3rd Qu.:1.0000 3rd Qu.:1.703e+10 3rd Qu.:1.703e+10
## Max. :1.0000 Max. :1.703e+10 Max. :1.703e+10
## NA's :1 NA's :775 NA's :775
## Pickup_Community_Area Dropoff_Community_Area Fare Tip
## Min. : 1.00 Min. : 1.00 Min. : 0.00 Min. : 0.000
## 1st Qu.: 8.00 1st Qu.: 8.00 1st Qu.: 10.00 1st Qu.: 0.000
## Median :25.00 Median :28.00 Median : 15.00 Median : 0.000
## Mean :28.14 Mean :29.01 Mean : 19.02 Mean : 1.527
## 3rd Qu.:37.00 3rd Qu.:38.00 3rd Qu.: 22.50 3rd Qu.: 2.000
## Max. :77.00 Max. :77.00 Max. :110.00 Max. :38.000
## NA's :189 NA's :185 NA's :1 NA's :1
## Additional_Charges Trip_Total Shared_Trip_Authorized Shared_Trip_Match
## Min. : 0.000 Min. : 0.00 Length:2000 Length:2000

```



```
## 1st Qu.: 2.010      1st Qu.: 13.73   Class :character      Class :character
## Median : 3.700      Median : 19.43   Mode  :character      Mode  :character
## Mean   : 4.691      Mean   : 25.24
## 3rd Qu.: 5.715      3rd Qu.: 29.69
## Max.   :62.680      Max.   :129.41
## NA's   :1          NA's    :1
## Trips_Pooled Pickup_Centroid_Latitude Pickup_Centroid_Longitude
## Min.      :1.000      Min.      :41.66      Min.      : -87.91
## 1st Qu.    :1.000      1st Qu.    :41.87      1st Qu.    : -87.69
## Median     :1.000      Median     :41.89      Median     : -87.65
## Mean       :1.042      Mean       :41.89      Mean       : -87.67
## 3rd Qu.    :1.000      3rd Qu.    :41.94      3rd Qu.    : -87.63
## Max.       :5.000      Max.       :42.02      Max.       : -87.53
##           NA's      :183      NA's      :183
## Pickup_Centroid_Location Dropoff_Centroid_Latitude Dropoff_Centroid_Longitude
## Length:2000      Min.      :41.66      Min.      : -87.91
## Class :character      1st Qu. :41.87      1st Qu.    : -87.70
## Mode  :character      Median   :41.89      Median     : -87.66
##           Mean       :41.89      Mean       : -87.67
##           3rd Qu.    :41.94      3rd Qu.    : -87.63
##           Max.       :42.02      Max.       : -87.55
##           NA's      :177      NA's      :177
## Dropoff_Centroid_Location
## Length:2000
## Class :character
## Mode  :character
##
##
##
##
```

Add some calc columns

```
trips$Trip_Start_Timestamp = as.Date(trips$Trip_Start_Timestamp, format="%m/%d/%Y")
trips$dow = weekdays(trips$Trip_Start_Timestamp)
head(trips)
```

```
##      PK_UID      Trip_ID Trip_Start_Timestamp
## 1 1078724 a45e404c67637a35ba8c68eb02abd0ef57a27851      2024-06-26
## 2 5948098 9bd17c2ea621b7c1bba3725d89c18b45ac574971      2024-06-07
## 3 1058729 92cc46f8531e33eee0e5a21cda3cb34d7997e0c8      2024-06-27
## 4 2375826 4c4659ad8d22f0f0177e79e2140023a9ffea9ea2      2024-06-21
## 5 559953 aceb9a135afb3a822bce97dbf5f80b31fa951c3d      2024-06-28
## 6 7199958 8fd92610f23c8513b2618355ea3dbaf1ac734e37      2024-06-02
##      Trip_End_Timestamp Trip_Seconds Trip_Miles Percent_Time_Chicago
## 1 06/26/2024 10:30:00 PM      1124      4.9      1
## 2 06/07/2024 03:30:00 PM      896      3.0      1
## 3 06/27/2024 01:30:00 AM      1994      27.8      0
## 4 06/21/2024 09:15:00 PM      2963      26.9      1
## 5 06/28/2024 10:45:00 PM      816      2.4      1
## 6 06/02/2024 06:00:00 AM      2260      32.4      1
##      Percent_Distance_Chicago Pickup_Census_Tract Dropoff_Census_Tract
## 1      1      NA      NA
## 2      1      NA      NA
## 3      0      17031980000      NA
## 4      1      NA      NA
## 5      1      17031081403      17031842200
## 6      1      17031820901      17031980000
##      Pickup_Community_Area Dropoff_Community_Area Fare Tip Additional_Charges
```

```

## 1      28      6 32.5  0      1.23
## 2      28      24 17.5  0      5.68
## 3      76      NA 60.0 12     21.39
## 4      NA      6 22.5  0      6.95
## 5      8       8 10.0  0      1.23
## 6      NA     76 47.5 12     28.12
## Trip_Total Shared_Trip_Authorized Shared_Trip_Match Trips_Pooled
## 1      33.73      false      false      1
## 2      23.18      false      false      1
## 3      93.39      false      false      1
## 4      29.45      false      false      1
## 5      11.23      false      false      1
## 6      87.62      false      false      1
## Pickup_Centroid_Latitude Pickup_Centroid_Longitude
## 1      41.87400      -87.66352
## 2      41.87400      -87.66352
## 3      41.97907      -87.90304
## 4      NA          NA
## 5      41.89092      -87.61887
## 6      NA          NA
## Pickup_Centroid_Location Dropoff_Centroid_Latitude
## 1 POINT (-87.6635175498 41.874005383)      41.94423
## 2 POINT (-87.6635175498 41.874005383)      41.90121
## 3 POINT (-87.9030396611 41.9790708201)      NA
## 4      41.94423
## 5 POINT (-87.6188683546 41.8909220259)      41.90494
## 6      41.97907
## Dropoff_Centroid_Longitude Dropoff_Centroid_Location dow
## 1      -87.65600 POINT (-87.6559981815 41.9442266014) Wednesday
## 2      -87.67636 POINT (-87.6763559892 41.9012069941) Friday
## 3      NA          Thursday
## 4      -87.65600 POINT (-87.6559981815 41.9442266014) Friday
## 5      -87.64991 POINT (-87.6499072264 41.9049353016) Friday
## 6      -87.90304 POINT (-87.9030396611 41.9790708201) Sunday

```

summary(trips)

```

## PK_UID Trip_ID Trip_Start_Timestamp Trip_End_Timestamp
## Min. : 2180 Length:2000 Min. :2024-06-01 Length:2000
## 1st Qu.:1857128 Class :character 1st Qu.:2024-06-08 Class :character
## Median :3728684 Mode :character Median :2024-06-15 Mode :character
## Mean :3749083 Mean :2024-06-15
## 3rd Qu.:5638863 3rd Qu.:2024-06-23
## Max. :7588816 Max. :2024-06-30
##
## Trip_Seconds Trip_Miles Percent_Time_Chicago
## Min. : 55.0 Min. : 0.200 Min. :0.0000
## 1st Qu.: 615.8 1st Qu.: 2.175 1st Qu.:1.0000
## Median : 991.5 Median : 4.400 Median :1.0000
## Mean :1218.3 Mean : 7.336 Mean :0.9325
## 3rd Qu.:1583.0 3rd Qu.:10.000 3rd Qu.:1.0000
## Max. :5937.0 Max. :63.900 Max. :1.0000
## NA's :1
## Percent_Distance_Chicago Pickup_Census_Tract Dropoff_Census_Tract
## Min. :0.0000 Min. :1.703e+10 Min. :1.703e+10
## 1st Qu.:1.0000 1st Qu.:1.703e+10 1st Qu.:1.703e+10
## Median :1.0000 Median :1.703e+10 Median :1.703e+10
## Mean :0.9275 Mean :1.703e+10 Mean :1.703e+10
## 3rd Qu.:1.0000 3rd Qu.:1.703e+10 3rd Qu.:1.703e+10
## Max. :1.0000 Max. :1.703e+10 Max. :1.703e+10

```

```

## NA's :1          NA's :775          NA's :775
## Pickup_Community_Area Dropoff_Community_Area Fare Tip
## Min. : 1.00      Min. : 1.00      Min. : 0.00 Min. : 0.000
## 1st Qu.: 8.00      1st Qu.: 8.00      1st Qu.: 10.00 1st Qu.: 0.000
## Median :25.00      Median :28.00      Median : 15.00 Median : 0.000
## Mean :28.14      Mean :29.01      Mean : 19.02 Mean : 1.527
## 3rd Qu.:37.00      3rd Qu.:38.00      3rd Qu.: 22.50 3rd Qu.: 2.000
## Max. :77.00      Max. :77.00      Max. :110.00 Max. :38.000
## NA's :189      NA's :185      NA's :1      NA's :1
## Additional_Charges Trip_Total Shared_Trip_Authorized Shared_Trip_Match
## Min. : 0.000      Min. : 0.00      Length:2000      Length:2000
## 1st Qu.: 2.010      1st Qu.: 13.73      Class :character      Class :character
## Median : 3.700      Median : 19.43      Mode :character      Mode :character
## Mean : 4.691      Mean : 25.24
## 3rd Qu.: 5.715      3rd Qu.: 29.69
## Max. :62.680      Max. :129.41
## NA's :1      NA's :1
## Trips_Pooled Pickup_Centroid_Latitude Pickup_Centroid_Longitude
## Min. :1.000      Min. :41.66      Min. : -87.91
## 1st Qu.:1.000      1st Qu.:41.87      1st Qu.: -87.69
## Median :1.000      Median :41.89      Median : -87.65
## Mean :1.042      Mean :41.89      Mean : -87.67
## 3rd Qu.:1.000      3rd Qu.:41.94      3rd Qu.: -87.63
## Max. :5.000      Max. :42.02      Max. : -87.53
## NA's :183      NA's :183
## Pickup_Centroid_Location Dropoff_Centroid_Latitude Dropoff_Centroid_Longitude
## Length:2000      Min. :41.66      Min. : -87.91
## Class :character      1st Qu.:41.87      1st Qu.: -87.70
## Mode :character      Median :41.89      Median : -87.66
## Mean :41.89      Mean : -87.67
## 3rd Qu.:41.94      3rd Qu.: -87.63
## Max. :42.02      Max. : -87.55
## NA's :177      NA's :177
## Dropoff_Centroid_Location dow
## Length:2000      Length:2000
## Class :character      Class :character
## Mode :character      Mode :character
##
##
##
##

```

Get rid of some not useful columns that contain info like identifiers and coordinates

```

trips_int = trips[,c("Trip_Seconds","Trip_Miles","Percent_Time_Chicago","Percent_Distance_Chicago","Fare","Tip","Additional_Charges")]
summary(trips_int)

```

```

## Trip_Seconds Trip_Miles Percent_Time_Chicago
## Min. : 55.0 Min. : 0.200 Min. :0.0000
## 1st Qu.: 615.8 1st Qu.: 2.175 1st Qu.:1.0000
## Median : 991.5 Median : 4.400 Median :1.0000
## Mean :1218.3 Mean : 7.336 Mean :0.9325
## 3rd Qu.:1583.0 3rd Qu.:10.000 3rd Qu.:1.0000
## Max. :5937.0 Max. :63.900 Max. :1.0000
## NA's :1
## Percent_Distance_Chicago Fare Tip Additional_Charges
## Min. :0.0000 Min. : 0.00 Min. : 0.000 Min. : 0.000
## 1st Qu.:1.0000 1st Qu.: 10.00 1st Qu.: 0.000 1st Qu.: 2.010
## Median :1.0000 Median : 15.00 Median : 0.000 Median : 3.700

```

```
## Mean :0.9275          Mean : 19.02    Mean : 1.527    Mean : 4.691
## 3rd Qu.:1.0000        3rd Qu.: 22.50    3rd Qu.: 2.000    3rd Qu.: 5.715
## Max. :1.0000          Max. :110.00    Max. :38.000    Max. :62.680
## NA's :1              NA's :1        NA's :1        NA's :1
## Trip_Total    Trips_Pooled    Shared_Trip_Authorized    Shared_Trip_Match
## Min. : 0.00    Min. :1.000    Length:2000            Length:2000
## 1st Qu.: 13.73    1st Qu.:1.000    Class :character        Class :character
## Median : 19.43    Median :1.000    Mode :character         Mode :character
## Mean : 25.24    Mean :1.042
## 3rd Qu.: 29.69    3rd Qu.:1.000
## Max. :129.41    Max. :5.000
## NA's :1
## dow
## Length:2000
## Class :character
## Mode :character
##
##
##
```

Check if the trip_total column is just the sum of the other three cost columns and get rid of it if so to not cause dependent columns

```
trips_int$total_check = with(trips_int, trips_int$Fare + trips_int$Tip + trips_int$Tip + trips_int$Additional_Charge)
summary(trips_int$total_check)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.    NA's
## 0.000   0.000   0.000   1.527   2.000   38.000     1
```

Its inconsitent so we will clean by just dropping this column to be safe

```
trips_int$total_check = NULL
trips_int$Trip_Total = NULL
```

Check for nulls

```
head(trips_int)
```

```
## Trip_Seconds Trip_Miles Percent_Time_Chicago Percent_Distance_Chicago Fare
## 1      1124      4.9              1              1 32.5
## 2      896      3.0              1              1 17.5
## 3     1994     27.8              0              0 60.0
## 4     2963     26.9              1              1 22.5
## 5      816      2.4              1              1 10.0
## 6     2260     32.4              1              1 47.5
## Tip Additional_Charges Trips_Pooled Shared_Trip_Authorized Shared_Trip_Match
## 1  0              1.23          1              false          false
## 2  0              5.68          1              false          false
## 3 12             21.39          1              false          false
## 4  0              6.95          1              false          false
## 5  0              1.23          1              false          false
## 6 12             28.12          1              false          false
##      dow
## 1 Wednesday
## 2   Friday
## 3 Thursday
## 4   Friday
```

```
sum(is.na(trips_int))
```

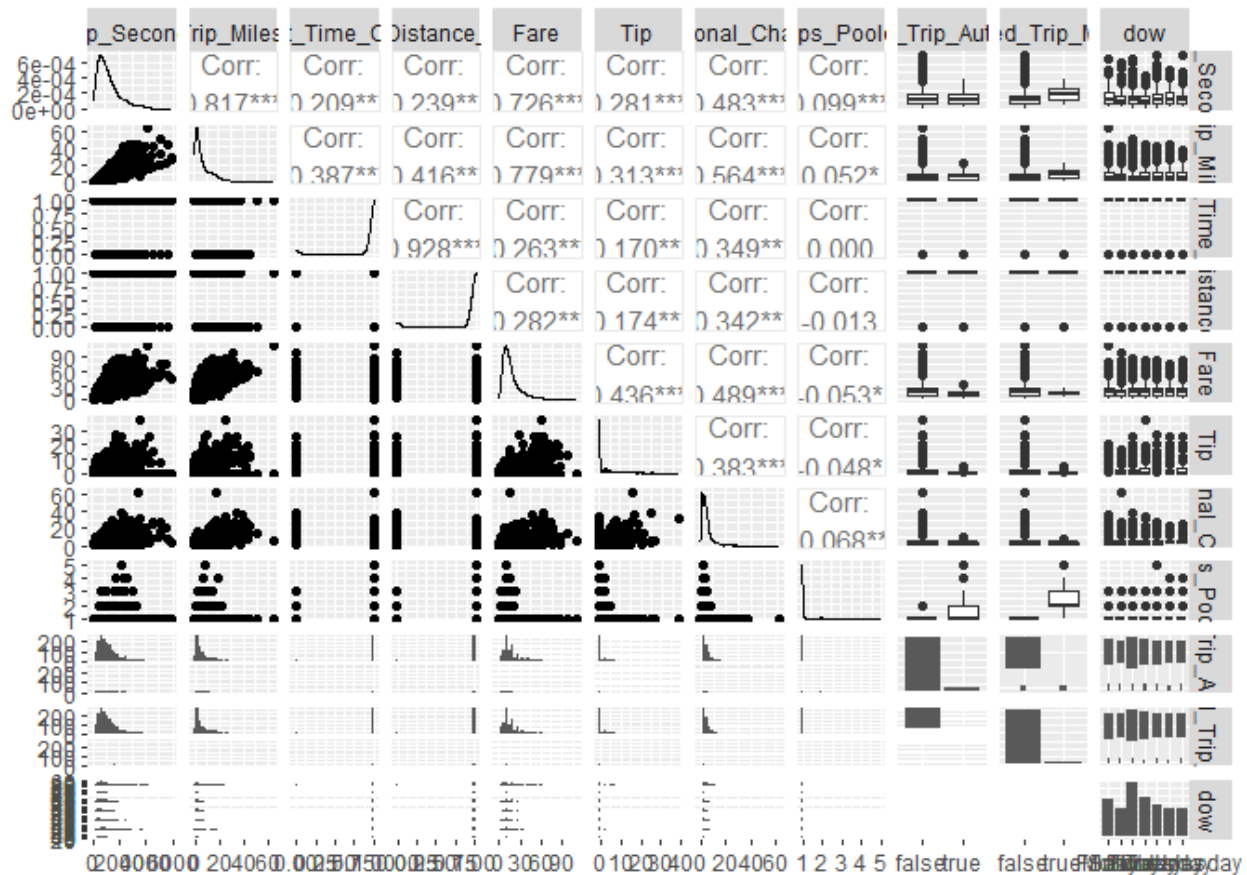
```
trips_int=trips_int[complete.cases(trips_int),]
nrow(trips_int)
```

```
## [1] 1998
```

```
set.seed(1)
train_ind = sample(1:nrow(trips_int),round(nrow(trips_int)*0.9))
train <- trips_int[train_ind,]
test <- trips_int[-train_ind,]
```

```
ggpairs(train)
```

[illegible]



the two percent columns also do not seem useful so lets get rid of them too.

```
trips_int$Percent_Time_Chicago = NULL
trips_int$Percent_Distance_Chicago = NULL
train <- trips_int[train_ind,]
test <- trips_int[-train_ind,]
```

I want to create a fare model from this data to use in my research. Lets find out what variables are important to fare.

```
mod = lm(Fare~.,data = train)
summary(mod)
```

```
##
## Call:
## lm(formula = Fare ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -31.136  -4.114  -1.440   2.645  52.441
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   6.894924   1.621669   4.252 2.23e-05 ***
## Trip_Seconds    0.004439   0.000379  11.714 < 2e-16 ***
## Trip_Miles     0.930867   0.045233  20.579 < 2e-16 ***
## Tip            0.798935   0.059870  13.344 < 2e-16 ***
```

```
## Additional_Charges      -0.030049   0.053778  -0.559   0.5764
## Trips_Pooled            -0.792071   1.531302  -0.517   0.6050
## Shared_Trip_Authorizedtrue -4.890068   1.067234  -4.582  4.92e-06 ***
## Shared_Trip_Matchtrue   -2.971141   2.640581  -1.125   0.2607
## dowMonday               -0.085948   0.729787  -0.118   0.9063
## dowSaturday              0.458279   0.623896   0.735   0.4627
## dowSunday                1.223020   0.671381   1.822   0.0687 .
## dowThursday              -0.705050   0.703864  -1.002   0.3166
## dowTuesday               -0.613870   0.728873  -0.842   0.3998
## dowWednesday             -0.520356   0.725991  -0.717   0.4736
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.886 on 1784 degrees of freedom
## Multiple R-squared:  0.6857, Adjusted R-squared:  0.6835
## F-statistic: 299.5 on 13 and 1784 DF, p-value: < 2.2e-16
```

DOW doesn't appear to affect things but tip does, and thats of concern because there is an intuitive relationship between tip and fare. So we should drop this. But lets check correlation to be sure.

```
train_num = subset(train, select = -c(dow,Shared_Trip_Authorized,Shared_Trip_Match))
test_num = subset(test, select = -c(dow,Shared_Trip_Authorized,Shared_Trip_Match))

cor_mat = cor(train_num)
cor_mat
```

```
##          Trip_Seconds Trip_Miles      Fare      Tip
## Trip_Seconds      1.00000000  0.8167555  0.72560059  0.28135121
## Trip_Miles        0.81675553  1.00000000  0.77914316  0.31294200
## Fare              0.72560059  0.7791432  1.00000000  0.43593463
## Tip               0.28135121  0.3129420  0.43593463  1.00000000
## Additional_Charges 0.48296814  0.5638116  0.48942965  0.38281205
## Trips_Pooled       0.09920895  0.0524975 -0.05279474 -0.04785573
##          Additional_Charges Trips_Pooled
## Trip_Seconds      0.48296814  0.09920895
## Trip_Miles        0.56381163  0.05249750
## Fare              0.48942965 -0.05279474
## Tip               0.38281205 -0.04785573
## Additional_Charges 1.00000000 -0.06813661
## Trips_Pooled      -0.06813661  1.00000000
```

Trip miles and trip seconds are highly correlated but thats fine, most fare models use both so we will hold onto it for now.

```
trips_int$Tip = NULL
train <- trips_int[train_ind,]
test <- trips_int[-train_ind,]
mod_lm = lm(Fare~.,data=train)
summary(mod_lm)
```

```
##
## Call:
## lm(formula = Fare ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -31.048  -4.116  -1.617   2.567  53.791
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    6.5729591   1.7000144   3.866 0.000114 ***
## Trip_Seconds    0.0046652   0.0003969  11.754 < 2e-16 ***
## Trip_Miles      0.9636633   0.0473538  20.350 < 2e-16 ***
```

```
## Additional_Charges      0.1529081  0.0545190   2.805 0.005091 **
## Trips_Pooled           -0.7980036  1.6054593  -0.497 0.619211
## Shared_Trip_Authorizedtrue -5.4702694  1.1179887  -4.893 1.08e-06 ***
## Shared_Trip_Matchtrue   -3.1061418  2.7684382  -1.122 0.262020
## dowMonday              0.1151217  0.7649658   0.150 0.880393
## dowSaturday            0.7886384  0.6535943   1.207 0.227739
## dowSunday              1.6976677  0.7029062   2.415 0.015826 *
## dowThursday            -0.7321640  0.7379478  -0.992 0.321253
## dowTuesday             -0.4810313  0.7640993  -0.630 0.529076
## dowWednesday           -0.2442809  0.7608396  -0.321 0.748197
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.268 on 1785 degrees of freedom
## Multiple R-squared:  0.6544, Adjusted R-squared:  0.6521
## F-statistic: 281.6 on 12 and 1785 DF, p-value: < 2.2e-16
```

LM

Only Shared_Trip_Authorized, trip_miles, and trip_seconds were significant so lets use these only. Additional_Charges is significant but this is not a useful variable for our use case because this is not something we could know ahead of time or it would probably be fixed (taxes and fees). Sunday is now slightly significant so we will include that too.

```
mod_lm = lm(Fare~dow+Shared_Trip_Authorized+Trip_Miles+Trip_Seconds,data=train)
summary(mod_lm)
```

```
##
## Call:
## lm(formula = Fare ~ dow + Shared_Trip_Authorized + Trip_Miles +
##     Trip_Seconds, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -27.049  -4.250  -1.698   2.665  54.253
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    6.2068590   0.5934005   10.460  <2e-16 ***
## dowMonday      0.2200200   0.7669919    0.287   0.774
## dowSaturday    0.7154411   0.6557313    1.091   0.275
## dowSunday      1.7009283   0.7052920    2.412   0.016 *
## dowThursday   -0.6983884   0.7400216   -0.944   0.345
## dowTuesday    -0.4207460   0.7658362   -0.549   0.583
## dowWednesday  -0.2132377   0.7631928   -0.279   0.780
## Shared_Trip_Authorizedtrue -7.7497278   0.8284607  -9.354  <2e-16 ***
## Trip_Miles     1.0137258   0.0447858  22.635  <2e-16 ***
## Trip_Seconds   0.0046016   0.0003953  11.640  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.299 on 1788 degrees of freedom
## Multiple R-squared:  0.6512, Adjusted R-squared:  0.6495
## F-statistic: 370.9 on 9 and 1788 DF, p-value: < 2.2e-16
```

Simple LM has r^2 of only 0.65, lets check its predictive accuracy.

```
test$lm_preds = predict(mod_lm,newdata = test)
mse <- mean((test$Fare - test$lm_preds)^2)
paste("MSE=",mse)
```

```
## [1] "MSE= 52.2219365866465"
```


Lets use subsets to check and see if anything else is helpful

```
regfit_full = regsubsets(Fare~.,train)
summary(regfit_full)

## Subset selection object
## Call: regsubsets.formula(Fare ~ ., train)
## 12 Variables (and intercept)
##
```

		Forced in	Forced out
## Trip_Seconds		FALSE	FALSE
## Trip_Miles		FALSE	FALSE
## Additional_Charges		FALSE	FALSE
## Trips_Pooled		FALSE	FALSE
## Shared_Trip_Authorizedtrue		FALSE	FALSE
## Shared_Trip_Matchtrue		FALSE	FALSE
## dowMonday		FALSE	FALSE
## dowSaturday		FALSE	FALSE
## dowSunday		FALSE	FALSE
## dowThursday		FALSE	FALSE
## dowTuesday		FALSE	FALSE
## dowWednesday		FALSE	FALSE

```
## 1 subsets of each size up to 8
## Selection Algorithm: exhaustive
##
```

		Trip_Seconds	Trip_Miles	Additional_Charges	Trips_Pooled
## 1 (1)	" "	"*"	" "	" "	" "
## 2 (1)	"*"	"*"	" "	" "	" "
## 3 (1)	"*"	"*"	" "	" "	" "
## 4 (1)	"*"	"*"	" "	" "	" "
## 5 (1)	"*"	"*"	"*"	" "	" "
## 6 (1)	"*"	"*"	"*"	" "	" "
## 7 (1)	"*"	"*"	"*"	" "	" "
## 8 (1)	"*"	"*"	"*"	" "	" "

```
##
```

		Shared_Trip_Authorizedtrue	Shared_Trip_Matchtrue	dowMonday	dowSaturday
## 1 (1)	" "	" "	" "	" "	" "
## 2 (1)	" "	" "	" "	" "	" "
## 3 (1)	"*"	" "	" "	" "	" "
## 4 (1)	"*"	" "	" "	" "	" "
## 5 (1)	"*"	" "	" "	" "	" "
## 6 (1)	"*"	"*"	" "	" "	" "
## 7 (1)	"*"	"*"	" "	" "	"*"
## 8 (1)	"*"	"*"	" "	" "	"*"

```
##
```

		dowSunday	dowThursday	dowTuesday	dowWednesday
## 1 (1)	" "	" "	" "	" "	" "
## 2 (1)	" "	" "	" "	" "	" "
## 3 (1)	" "	" "	" "	" "	" "
## 4 (1)	"*"	" "	" "	" "	" "
## 5 (1)	"*"	" "	" "	" "	" "
## 6 (1)	"*"	" "	" "	" "	" "
## 7 (1)	"*"	" "	" "	" "	" "
## 8 (1)	"*"	"*"	" "	" "	" "

```
regfit_full = regsubsets(Fare~.,data = train, nvmax=8)
reg_sum = summary(regfit_full)
names(reg_sum)

## [1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"

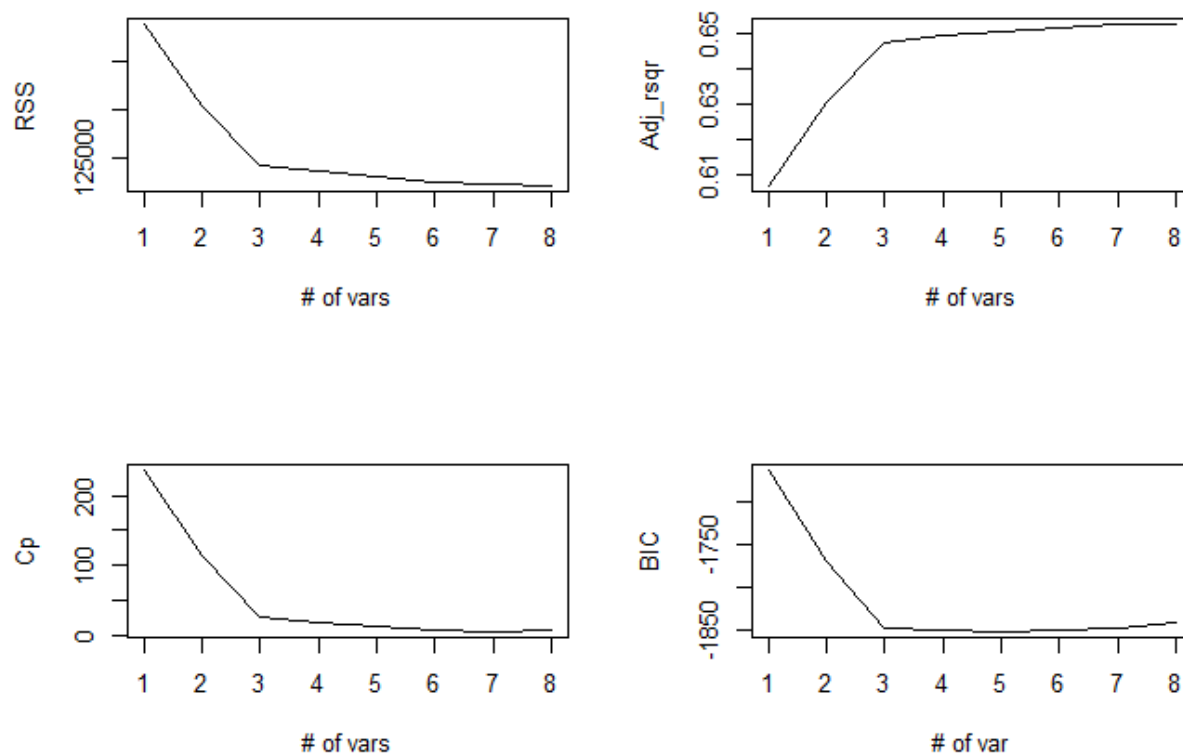
reg_sum$rsq

## [1] 0.6070641 0.6309810 0.6482925 0.6501786 0.6517873 0.6531103 0.6540107
## [8] 0.6541979
```

```

par(mfrow=c(2,2))
plot(reg_sum$rss,xlab = "# of vars",ylab="RSS",type="l")
plot(reg_sum$adjr2,xlab = "# of vars",ylab = "Adj_rsqr",type="l")
plot(reg_sum$cp,xlab="# of vars",ylab="Cp",type="l")
plot(reg_sum$bic, xlab="# of var",ylab="BIC",type="l")

```



These results are showing that some of the higher order models are better fits like 5 in BIC, 7 in Cp but the most obvious change is with the 3 variable model. Lets try stepwise selection.

```

null <-lm(Fare ~ 1, data=train)
full <- lm(Fare ~ ., data=train)

stepAIC(full, scope = list(lower = null, upper= full),direction = "both", trace = FALSE)

```

```

##
## Call:
## lm(formula = Fare ~ Trip_Seconds + Trip_Miles + Additional_Charges +
##     Shared_Trip_Authorized + Shared_Trip_Match + dow, data = train)
##
## Coefficients:
##              (Intercept)              Trip_Seconds
##                5.782927                0.004662
##              Trip_Miles              Additional_Charges
##                0.963446                0.153763
## Shared_Trip_Authorizedtrue Shared_Trip_Matchtrue
##               -5.476556               -4.232899
##              dowMonday              dowSaturday
##                0.119332                0.784337

```

```
##                dowSunday                dowThursday
##                1.689476                -0.745963
##                dowTuesday                dowWednesday
##                -0.498965                -0.255236


predict.regsbsets =function (object ,newdata ,id ,...){
  form=as.formula (object$call [[2]])
  mat=model.matrix (form ,newdata )
  coefi =coef(object ,id=id)
  xvars =names (coefi )
  mat[,xvars ]%*% coefi
}

k= 10
folds=sample(1:k,nrow(train),replace = TRUE)
cv.errors=matrix(NA,k,11,dimnames=list(NULL,paste(1:11)))
for (j in 1:k){
  best.fit=regsubsets(Fare~.,data=train[folds!=j,],nvmax=11)
  for (i in 1:11){
    pred=predict(best.fit,train[folds==j,],id=i)
    cv.errors[j,i]=mean((train$Fare[folds==j]-pred)^2)
  }
}
mean.cv.errors=apply(cv.errors,2,mean)
mean.cv.errors
```

```
##          1          2          3          4          5          6          7          8
## 77.23545 72.64004 69.28467 69.42865 69.30912 69.03773 68.52978 68.59183
##          9         10         11
## 68.73178 68.74530 68.73758
```

```
par(mfrow=c(1,1))
plot(mean.cv.errors,type="b")
which.min(mean.cv.errors)
```

```
## 7
## 7
points(7,mean.cv.errors[7],col="red",cex=2,pch=20)
```



The 8 variable model has the best mean.cv.errors.

```
reg.best = regsubsets(Fare~.,train,nvmax=11)
coef(reg.best,7)
```

```
##          (Intercept)          Trip_Seconds
##          5.521221404          0.004651536
##          Trip_Miles          Additional_Charges
##          0.963793543          0.154063452
## Shared_Trip_Authorizedtrue Shared_Trip_Matchtrue
```


```
##           -5.465101328           -4.217494625
##           dowSaturday           dowSunday
##           1.053798586           1.958618473
test$predicted_y_cv_subset <- predict.regsbsets(reg.best, test, 7)
mse <- mean((test$Fare - test$predicted_y_cv_subset)^2)
mse
```

```
## [1] 51.64421
ss_total <- sum((test$Fare - mean(test$Fare))^2)
ss_res <- sum((test$Fare - test$predicted_y)^2)
r_squared <- 1 - (ss_res / ss_total)
r_squared
```

```
## [1] 0.6314024
```

R^2 is still worse than LM, but MSE is slightly better. Less bias in the cv subset model.

```
x=model.matrix(Fare~.,train_num)[,-1]
y=train_num$Fare
grid = 10^seq(10,-2,length=100)
ridge.mod = glmnet(x,y, alpha = 0,lambda=grid)
cv.out.ridge = cv.glmnet(x,y,alpha=0)
plot(cv.out.ridge)
```



```
bestlam = cv.out.ridge$lambda.min
bestlam
```

```
## [1] 1.091826
```

```
test_x =model.matrix(Fare~.,test_num)[-1]
test_x
```

```
##      Trip_Seconds Trip_Miles Tip Additional_Charges Trips_Pooled
## 6              2260       32.4  12              28.12              1
```

## 10	421	1.3	0	1.23	1
## 24	908	3.6	2	1.18	2
## 25	592	1.4	0	2.94	1
## 30	541	1.8	0	1.23	1
## 32	1342	9.1	0	3.30	1
## 70	300	0.9	0	2.49	1
## 85	1379	8.1	0	7.02	1
## 123	980	2.2	3	4.46	1
## 154	919	6.6	0	5.35	1
## 155	2254	9.3	0	4.71	1
## 172	1024	5.1	3	3.45	1
## 178	482	4.1	0	3.89	1
## 188	2690	13.5	0	4.49	2
## 189	924	3.5	0	4.87	1
## 195	1072	6.0	6	7.90	1
## 200	444	0.7	3	4.89	1
## 211	1300	4.9	2	1.23	1
## 215	756	2.1	2	2.78	1
## 257	1273	8.2	10	2.98	1
## 263	783	2.5	2	1.23	1
## 283	943	10.6	0	6.83	1
## 288	908	4.7	0	1.23	1
## 295	366	1.1	3	3.77	1
## 297	309	1.3	0	1.23	1
## 301	459	1.5	3	2.01	1
## 331	2069	10.2	6	12.23	1
## 337	371	0.7	0	2.07	1
## 344	1868	13.0	14	1.23	1
## 362	697	2.3	0	2.52	1
## 367	731	4.2	0	3.29	1
## 370	1269	16.5	0	10.77	1
## 380	1421	5.6	0	2.98	1
## 385	1344	6.9	0	2.82	1
## 395	532	1.3	3	3.76	1
## 424	521	1.5	0	2.15	1
## 429	1652	7.2	0	6.00	1
## 433	1828	18.0	0	7.98	1
## 447	2204	13.2	0	4.57	1
## 475	1289	10.2	0	7.10	1
## 497	1973	9.0	0	1.23	1
## 523	97	0.4	0	2.68	1
## 527	527	1.3	3	4.35	1
## 531	759	3.0	0	1.23	1
## 542	492	2.0	0	4.27	1
## 545	1197	4.4	0	3.08	1
## 555	413	3.0	0	3.70	1
## 558	1814	8.6	0	5.89	1
## 602	681	3.2	0	3.93	1
## 603	1498	11.1	0	1.23	1
## 605	342	1.4	0	3.70	1
## 613	646	3.4	0	2.97	1
## 622	421	2.6	0	1.23	1
## 667	976	6.6	0	5.08	1
## 668	324	1.3	0	1.98	1
## 695	1625	17.0	9	11.23	1
## 697	1599	5.0	0	4.50	1
## 698	2512	14.4	5	6.23	1
## 700	1517	16.0	0	1.23	1
## 701	421	2.0	0	2.61	1

## 716	1154	9.8	0	1.23	1
## 734	1697	9.6	0	2.98	1
## 735	3107	16.4	0	5.55	1
## 742	283	1.1	0	1.88	1
## 755	288	0.8	3	2.16	1
## 772	755	6.7	3	3.08	1
## 803	353	1.2	0	2.82	1
## 806	1643	22.5	0	7.13	1
## 807	3205	14.8	0	9.98	1
## 814	2759	34.6	0	6.45	1
## 830	605	3.8	0	6.52	1
## 833	2454	9.9	0	4.52	2
## 834	788	3.4	0	3.10	1
## 870	1050	3.8	0	1.23	1
## 872	1911	16.5	0	7.06	1
## 884	459	2.8	0	3.32	1
## 888	671	2.0	1	4.46	1
## 893	1082	12.2	10	22.07	1
## 926	573	2.8	0	3.96	1
## 929	689	2.7	0	5.06	1
## 933	658	3.0	0	3.83	1
## 943	338	0.8	0	2.49	1
## 962	982	3.6	0	5.42	1
## 963	1171	9.4	0	4.81	1
## 964	1779	6.9	0	5.34	2
## 965	339	0.7	1	2.50	1
## 973	2650	13.7	0	5.56	1
## 982	476	1.8	1	2.73	1
## 1000	1836	5.6	0	6.14	1
## 1001	289	1.0	0	1.23	1
## 1003	3057	37.4	7	10.50	1
## 1005	1629	10.0	0	8.76	1
## 1034	1544	11.6	0	5.95	1
## 1044	1162	4.7	0	2.71	1
## 1058	3241	36.5	0	4.75	1
## 1063	778	2.4	0	3.45	1
## 1070	904	2.2	0	1.23	1
## 1071	319	1.0	0	8.10	1
## 1074	712	4.4	0	2.90	1
## 1077	370	0.8	5	1.23	1
## 1083	943	3.5	0	1.94	2
## 1086	1670	6.4	8	6.17	1
## 1097	357	1.6	0	1.84	1
## 1100	1070	7.1	0	5.48	1
## 1102	1768	12.9	0	6.11	1
## 1106	1779	8.4	0	4.17	1
## 1130	1089	5.5	0	7.08	1
## 1132	2254	13.4	7	5.46	1
## 1138	1007	8.2	4	1.23	1
## 1160	1148	13.9	0	19.64	1
## 1164	271	1.3	0	1.73	1
## 1173	771	2.1	0	4.23	1
## 1175	744	2.6	0	1.23	1
## 1185	1562	18.0	0	1.23	1
## 1193	804	7.4	0	3.34	1
## 1205	1702	9.4	0	5.06	1
## 1212	991	2.8	0	5.50	1
## 1223	1107	4.5	6	5.58	1
## 1227	292	1.3	0	2.78	1

## 1240	3985	14.5	20	6.23	1
## 1241	1985	8.3	0	6.75	1
## 1245	147	0.6	0	3.62	1
## 1248	641	1.8	2	1.23	1
## 1249	704	1.9	3	4.26	1
## 1263	870	2.3	0	2.98	1
## 1278	668	1.5	3	4.94	1
## 1283	510	1.4	0	2.98	1
## 1293	2251	19.2	0	6.23	1
## 1307	880	3.3	3	1.23	1
## 1311	704	2.2	0	1.23	1
## 1315	757	2.2	0	2.98	1
## 1321	1977	16.3	5	5.49	1
## 1322	671	1.4	0	3.82	1
## 1325	985	4.4	0	1.23	1
## 1335	908	7.2	0	5.70	1
## 1337	719	4.6	0	3.31	1
## 1339	995	5.1	0	1.23	1
## 1357	1815	17.2	7	22.26	1
## 1397	586	2.0	0	1.23	1
## 1409	426	1.5	2	1.23	1
## 1413	719	4.3	0	4.57	1
## 1415	484	1.0	3	2.65	1
## 1418	3608	33.3	0	8.47	1
## 1429	4337	40.4	5	7.23	1
## 1430	1549	14.8	0	6.15	1
## 1449	1286	11.7	0	1.23	1
## 1455	2164	31.9	8	5.72	1
## 1467	592	3.5	0	6.23	1
## 1476	814	4.8	0	1.23	1
## 1484	496	4.3	0	1.23	1
## 1489	632	3.7	0	4.11	1
## 1493	2087	9.0	0	8.04	1
## 1499	1679	12.4	0	3.93	2
## 1502	396	1.6	0	2.89	1
## 1505	734	1.1	1	4.24	1
## 1515	1216	3.3	0	5.14	1
## 1519	773	3.0	3	4.82	1
## 1520	1637	15.6	7	17.37	1
## 1548	729	3.7	0	2.98	1
## 1564	1008	5.6	0	2.73	1
## 1588	774	2.7	0	1.23	1
## 1613	1632	10.3	0	5.08	1
## 1623	1421	13.0	4	1.23	1
## 1644	1787	4.2	0	1.91	2
## 1651	1664	15.4	5	3.74	1
## 1674	831	7.4	0	2.66	1
## 1699	1390	7.7	4	5.36	1
## 1703	1315	16.8	0	16.98	1
## 1710	3173	21.3	9	13.99	1
## 1719	954	3.0	0	2.98	1
## 1726	2029	8.5	0	1.23	1
## 1729	1063	3.5	0	1.23	1
## 1743	2686	9.7	10	6.23	1
## 1766	2119	18.1	9	4.82	1
## 1767	316	1.5	0	9.24	1
## 1783	1274	2.3	0	2.98	1
## 1789	935	3.1	5	3.32	1
## 1790	757	2.3	0	14.67	1

```
## 1804      1745      13.0  0      1.96      3
## 1813      1054       4.9  0      1.54      2
## 1824      1309       9.7  4      1.23      1
## 1846       629       3.4  0      5.64      1
## 1859      1560      20.2  0      7.45      1
## 1863      2126      18.8  0      7.98      1
## 1869       690       4.7  0      4.60      1
## 1870      1175       6.7  0      7.99      1
## 1873      2083       5.2  3      2.98      1
## 1878       710       3.8  0      1.23      1
## 1902       455       1.3  3      2.89      1
## 1919      2631      13.9  7      5.24      1
## 1922      1027       3.2  3      5.80      1
## 1927       759       2.8  1      4.58      1
## 1928       660       2.6  3      4.12      1
## 1930       511       3.5  0      1.23      1
## 1935       436       1.4  0      2.98      1
## 1954       351       1.3  0      1.23      1
## 1985      1377       6.0  0      1.23      1
## 1992       912       4.9  0      1.23      1
## 1994       446       0.9  0      1.23      1
## 1996       375       1.2  5      2.27      1
```

```
ridge.pred=predict(ridge.mod,s=bestlam,newx=test_x)
mse_ridge <- mean((test_num$Fare - ridge.pred)^2)
paste("MSE=",mse_ridge)
```

```
## [1] "MSE= 44.5698834306074"
```


Very high MSE, not a lot of variables so do not really need ridge anyways.

PCR/pls

```
pcr.mod=pcr(Fare~.,data = train,scale=TRUE,validation="CV")
summary(pcr.mod)
```

```
## Data:      X dimension: 1798 12
## Y dimension: 1798 1
## Fit method: svdpc
## Number of components considered: 12
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           14.02   13.93   8.920   8.877   8.746   8.749   8.749
## adjCV        14.02   13.94   8.914   8.871   8.731   8.740   8.745
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps
## CV           8.745   8.387   8.387   8.367   8.336   8.337
## adjCV        8.740   8.382   8.383   8.363   8.330   8.332
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X       20.6891  39.56   49.89   59.60   69.12   78.54   87.95   92.52
## Fare    0.9291  59.90   60.32   61.65   61.66   61.67   61.79   64.81
##      9 comps 10 comps 11 comps 12 comps
## X       96.19   97.73   99.14   100.00
## Fare    64.81   65.14   65.44   65.44
```

```
validationplot(pcr.mod,val.type="MSEP")
```



```
model_pcr_mse = MSEP(pcr.mod, estimate="CV")
model_pcr_mse
```

```
## (Intercept)      1 comps      2 comps      3 comps      4 comps      5 comps
##      196.59      194.05      79.56      78.80      76.49      76.55
##      6 comps      7 comps      8 comps      9 comps     10 comps     11 comps
##      76.54      76.47      70.33      70.35      70.00      69.48
##     12 comps
##      69.51
```

```

pcr.pred=predict(pcr.mod,test,ncomp=11)
mse_pcr <- mean((test$Fare - pcr.pred)^2)
paste("MSE=",mse_pcr)

```

```
## [1] "MSE= 52.0691020843782"
```

The model with 11 principal components performs the best.

```


pls.mod=plsr(Fare~.,data = train,scale=TRUE,validation="CV")
summary(pls.mod)

```

```

## Data:      X dimension: 1798 12
## Y dimension: 1798 1
## Fit method: kernelpls
## Number of components considered: 12
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           14.02   8.609   8.412   8.352   8.322   8.323   8.323
## adjCV         14.02   8.606   8.404   8.348   8.320   8.318   8.318
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps
## CV           8.324   8.324   8.324   8.324   8.324   8.324
## adjCV         8.319   8.319   8.319   8.320   8.320   8.320
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          18.78   26.78   39.89   53.32   55.84   65.56   69.02   76.53
## Fare       62.74   64.82   65.18   65.31   65.43   65.44   65.44   65.44
##      9 comps 10 comps 11 comps 12 comps
## X          80.15   84.69   90.57  100.00
## Fare       65.44   65.44   65.44   65.44
validationplot(pls.mod,val.type="MSEP")

```




```
model_pls_mse = MSEP(pls.mod, estimate="CV")
model_pls_mse
```

```
## (Intercept)      1 comps      2 comps      3 comps      4 comps      5 comps
##      196.59         74.12         70.76         69.75         69.26         69.27
##      6 comps      7 comps      8 comps      9 comps     10 comps     11 comps
##      69.27         69.28         69.29         69.29         69.29         69.29
##     12 comps
##      69.29
```

```
pls.pred=predict(pls.mod,test,ncomp=2)
mse_pls <- mean((test$Fare - pls.pred)^2)
paste("MSE=",mse_pls)
```

```
## [1] "MSE= 52.564886135774"
```

```
ncomp.onesigma <- selectNcomp(pls.mod, method = "onesigma", plot = TRUE)
```



Analyzing the msep validation plot, 2 principal components appear to be enough.

Final model

The LM was sufficient for my purposes so let's summarize it and interpret its coefficients. Remove DOW also since its relatively insignificant.

```
mod_lm_fin = lm(Fare~Shared_Trip_Authorized+Trip_Miles+Trip_Seconds,data=train)
summary(mod_lm_fin)
```

```
##
## Call:
## lm(formula = Fare ~ Shared_Trip_Authorized + Trip_Miles + Trip_Seconds,
##     data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -27.326  -4.232  -1.547   2.559  54.714
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      6.5967120   0.3491157   18.895  <2e-16 ***
## Shared_Trip_Authorizedtrue -7.7978116   0.8298213   -9.397  <2e-16 ***
## Trip_Miles         1.0212039   0.0448049   22.792  <2e-16 ***
## Trip_Seconds       0.0044736   0.0003942   11.349  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.32 on 1794 degrees of freedom
## Multiple R-squared:  0.6483, Adjusted R-squared:  0.6477
## F-statistic: 1102 on 3 and 1794 DF, p-value: < 2.2e-16

test$lm_preds_fin = predict(mod_lm_fin,newdata = test)
mse <- mean((test$Fare - test$lm_preds_fin)^2)
paste("MSE=",mse)

## [1] "MSE= 51.3003840981791"
```

References

- [1] Levy J. Transportation Network Providers-Trips (2018-2022)[supporting dataset]; 2023.
- [2] Gangadharaiah R, Su H, Rosopa EB, Brooks JO, Kolodge K, Boor L, et al. A User-Centered Design Exploration of Factors That Influence the Rideshare Experience. *Safety*. 2023;9. Available from: <https://www.mdpi.com/2313-576X/9/2/36>.
- [3] Gangadharaiah R, Brooks JO, Rosopa PJ, Su H, Boor L, Edgar A, et al. The Development of the Pooled Rideshare Acceptance Model (PRAM). *Safety*. 2023 9;9:61.
- [4] Su H, Gangadharaiah R, Rosopa EB, Brooks JO, Boor L, Kolodge K, et al. Exploration of Factors That Influence Willingness to Consider Pooled Rideshare. *Transportation Research Record: Journal of the Transportation Research Board*. 2024 1.
- [5] Auld J, Hope M, Ley H, Sokolov V, Xu B, Zhang K. POLARIS: Agent-based modeling framework development and implementation for integrated travel demand and network and operations simulations. *Transportation Research Part C: Emerging Technologies*. 2016;64:101-16. Available from: <https://www.sciencedirect.com/science/article/pii/S0968090X15002703>.
- [6] Paul J, Gurumurthy KM, Cokyasar T, Su H, Auld J, Jia Y. Optimization of Dynamic Ride-Sharing by Considering User Preference Through Discount and Delay Tolerance. In: 2023 IEEE 26th International Conference on Intelligent Transportation Systems (ITSC). IEEE; 2023. p. 2770-5.